

RAG Evaluation and Governance Platform

A local-first assurance project for testing retrieval quality, citation support, faithfulness, refusal behaviour and governance readiness.

Problem

Fluent RAG answers can still be unsupported or wrongly cited.

Method

Measure retrieval, citations, faithfulness, refusal and safety signals.

Result

Baseline E improves overall pass rate from 0.850 to 0.867.

Limit

Multi-hop pass rate remains weak at 0.200 on the main benchmark.

0.850

Baseline A

0.867

Baseline E

0.200

Multi-hop E

0.917

Citation E

Interpretation: retrieval improvement is useful but modest. Challenge-set results are separate from the main 60-question benchmark.